

NextRAN-AI – 03/10/2025

Integrated Systems Laboratory (ETH Zürich)

Marco Bertuletti

mbertuletti@iis.ee.ethz.ch

Yichao Zhang

yiczhang@iis.ee.ethz.ch

Mahdi Abdollahpour

mahdi.abdollahpour@unibo.it

Alessandro Vanelli-Coralli


avanelli@iis.ee.ethz.ch

Luca Benini

lbenini@iis.ee.ethz.ch

PULP Platform

Open Source Hardware, the way it should be!



pulp-platform.org

[@pulp_platform](https://twitter.com/pulp_platform)

[company/pulp-platform](https://www.linkedin.com/company/pulp-platform)

[youtube.com/pulp_platform](https://www.youtube.com/pulp_platform)



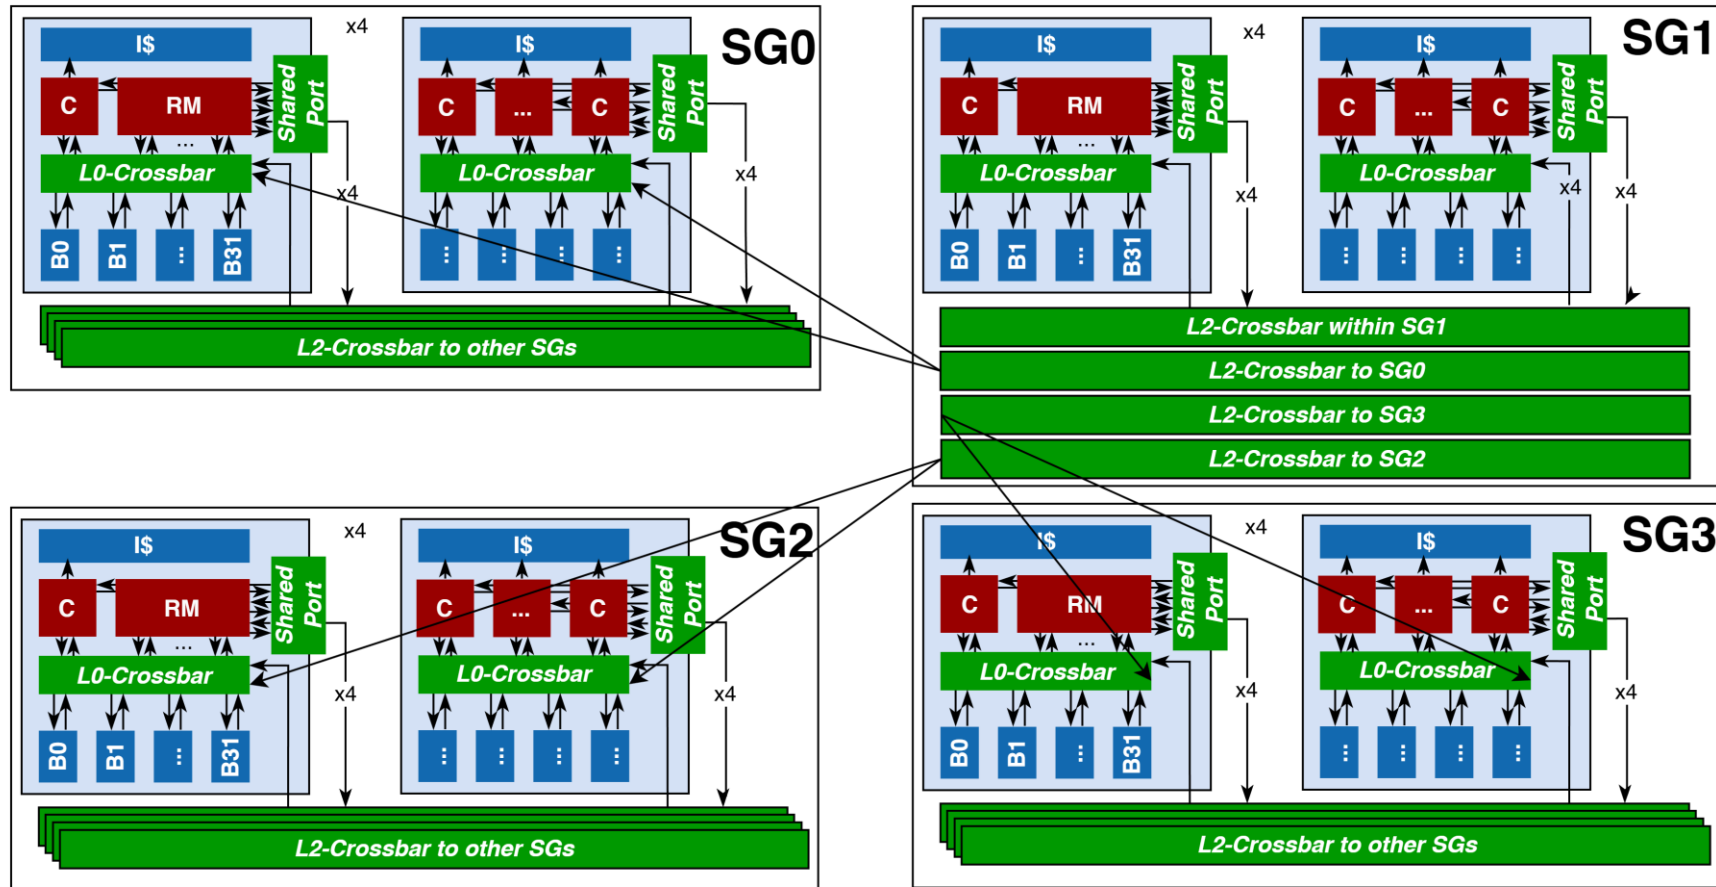
Outline

- Experiments on RedMulE size + Parallel RedMulE execution
- How to improve performance on large tensor arrays (16x16) → BW

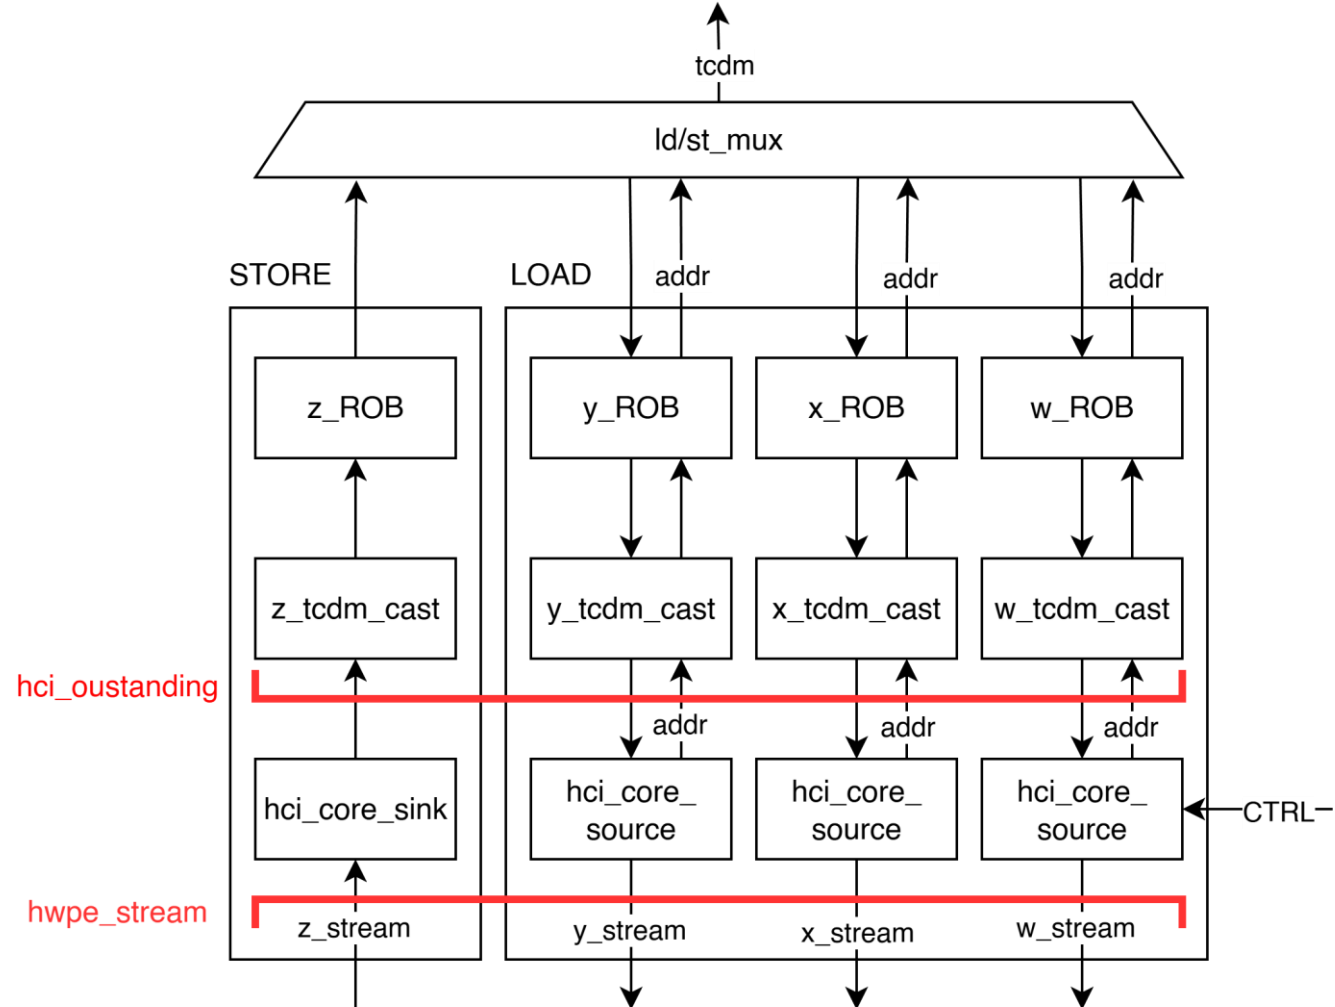


TensorPool SubGroup

100 cores, 4 RedMules, 4 Tiles, 8 cores / Tile, 128KiB SRAM



Implementation with three ROBs




Varying the size of the ROB




HWPE corresponds to the standalone RedMule

RedMule PEs Utilization



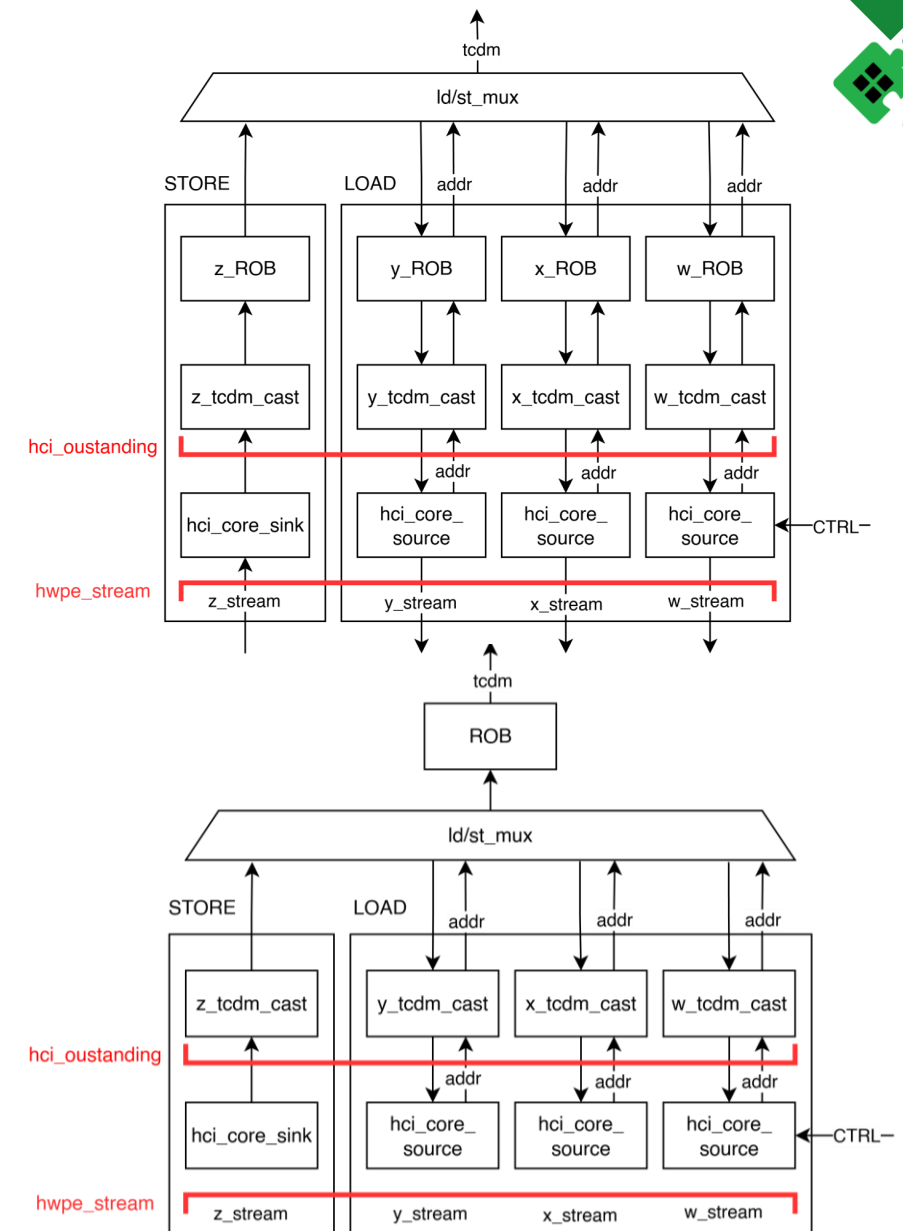
Runtime (cycles)



Implementation with one ROB


- Reduces size of the transaction table
- No waste in memory, because only some streams (W) need many ROB entries

ISSUE: transactions of different streams do not need to be in-order, if one stream is not ready it blocks the responses of other streams → **with dedicated ROBs we are 1.1x faster**



Parallelization

- Matmul parallelized over rows of X over multiple RedMulEs
- Modified the runtime to obtain easily redmule IDs and RedMulEs number



```
// Matmul, MxKxN, all data in L1
```


```
uint32_t redmule_id = get_redmule_id();
```

```
uint32_t redmule_num = get_redmule_num();
```

```
for (i = 0; i < num_redmules; i++) {  
    hwpe_soft_clear();  
    redmule_cfg (  
        ptr_X + redmule_id * M * K / num_redmules,  
        ptr_W,  
        ptr_Y + redmule_id * M * K / num_redmules,  
        M, K / num_redmules, N,  
        GEMM, Float16);  
    hwpe_trigger_job();  
}  
// potentially work from other cores  
mempool_barrier();
```



Parallelization Results



16x16 array requires 1024bit/cycle (32 memory ports)



- The accelerator memory ports access the same remote connection on writes and read responses
- To scale to 16x16 (256MACs/cycle) we should:
 - Increase the BW to other Tiles
 - Increase the L dimension of the accelerator, which does not impact the BW
 - Reduce the accelerator's BW



How to reduce the tensor engine BW?

We are solving a problem of size MKN, given the bitwidth of the elements β and tiling of size mnk we compute the bits loaded:

Output stationary:

```

1. for i:m:M
2.   for j:n:N
3.     load (Y)( $\beta \times m \times n$ )
4.     for l:k:K
5.       load (X)( $\beta \times m \times k$ )
6.       load (W)( $\beta \times k \times n$ )
7.       Y += X  $\times$  W
8.     end
9.     store (Y)( $\beta \times m \times n$ )
10.  end
11. end
  
```

$$\frac{M}{m} \frac{N}{n} \left(2\beta mn + \frac{K}{k} (\beta mk + \beta kn) \right) = MNK\beta \left(\frac{2}{K} + \frac{m+n}{mn} \right)$$

Input stationary:

```

1. for i:m:M
2.   for l:k:K
3.     load (X)( $\beta \times m \times k$ )
4.     for j:n:N
5.       load (Y)( $\beta \times m \times n$ )
6.       load (W)( $\beta \times k \times n$ )
7.       Y += X  $\times$  W
8.       store (Y)( $\beta \times m \times n$ )
9.     end
10.  end
11. end
  
```

$$\frac{M}{m} \frac{K}{k} \left(\beta mk + \frac{N}{n} (2\beta mn + \beta kn) \right) = MNK\beta \left(\frac{1}{N} + \frac{2m+k}{mk} \right)$$

Weight stationary:

```

1. for l:k:K
2.   for j:n:N
3.     load (W)( $\beta \times k \times n$ )
4.     for i:m:M
5.       load (Y)( $\beta \times m \times n$ )
6.       load (X)( $\beta \times m \times k$ )
7.       Y += X  $\times$  W
8.       store (Y)( $\beta \times m \times n$ )
9.     end
10.  end
11. end
  
```

$$\frac{N}{n} \frac{K}{k} \left(\beta kn + \frac{M}{m} (2\beta mn + \beta mk) \right) = MNK\beta \left(\frac{1}{M} + \frac{2n+k}{nk} \right)$$

How to reduce the tensor engine BW?




We compute the arithmetic intensity as the operations $2MKN$ divided by the bits loaded, therefore the bits loaded must be minimized:

- Output stationary $MNK\beta \left(\frac{2}{K} + \frac{m+n}{mn} \right) \approx \frac{MNK\beta(m+n)}{mn}$
- Input stationary $MNK\beta \left(\frac{1}{N} + \frac{2m+k}{mk} \right) \approx \frac{MNK\beta(2m+k)}{mk}$
- Weight stationary $MNK\beta \left(\frac{1}{M} + \frac{2n+k}{nk} \right) \approx \frac{MNK\beta(2n+k)}{nk}$

We notice that the smallest value of the loaded bits is for a square tile $m=n=k$, therefore the output stationary solution always guarantees the highest arithmetic intensity

Output stationary engine, outer product

- X and W inputs are broadcasted over row/columns
- Y inputs are stored in accumulators



x00 x01
x10 x11
x20 x21
x30 x31
...

w00 w01 w02 w03
w10 w11 w12 w13
...

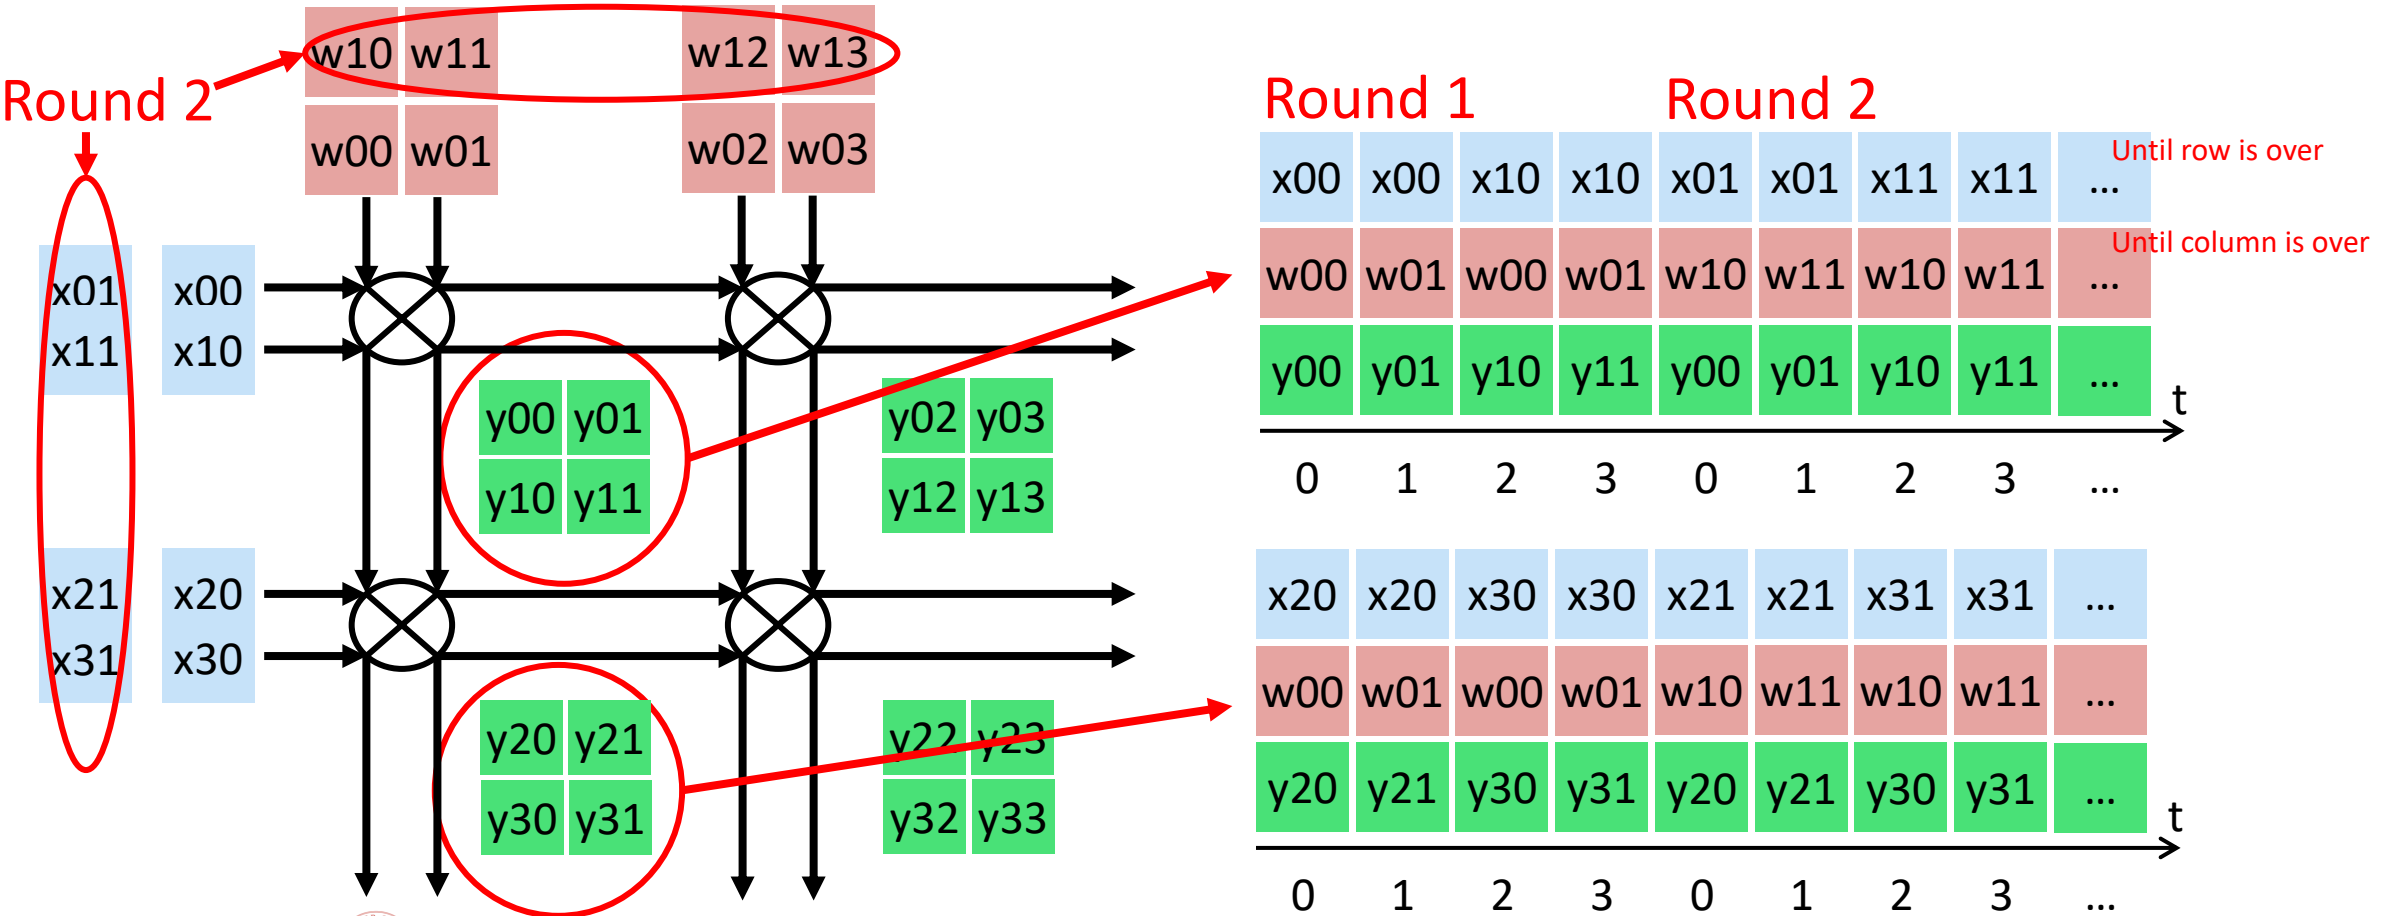
y00 y01 y02 y03
y10 y11 y12 y13
y20 y21 y22 y23
y30 y31 y32 y33
...

Rows in X are multiplied by
columns in W and stored in Y

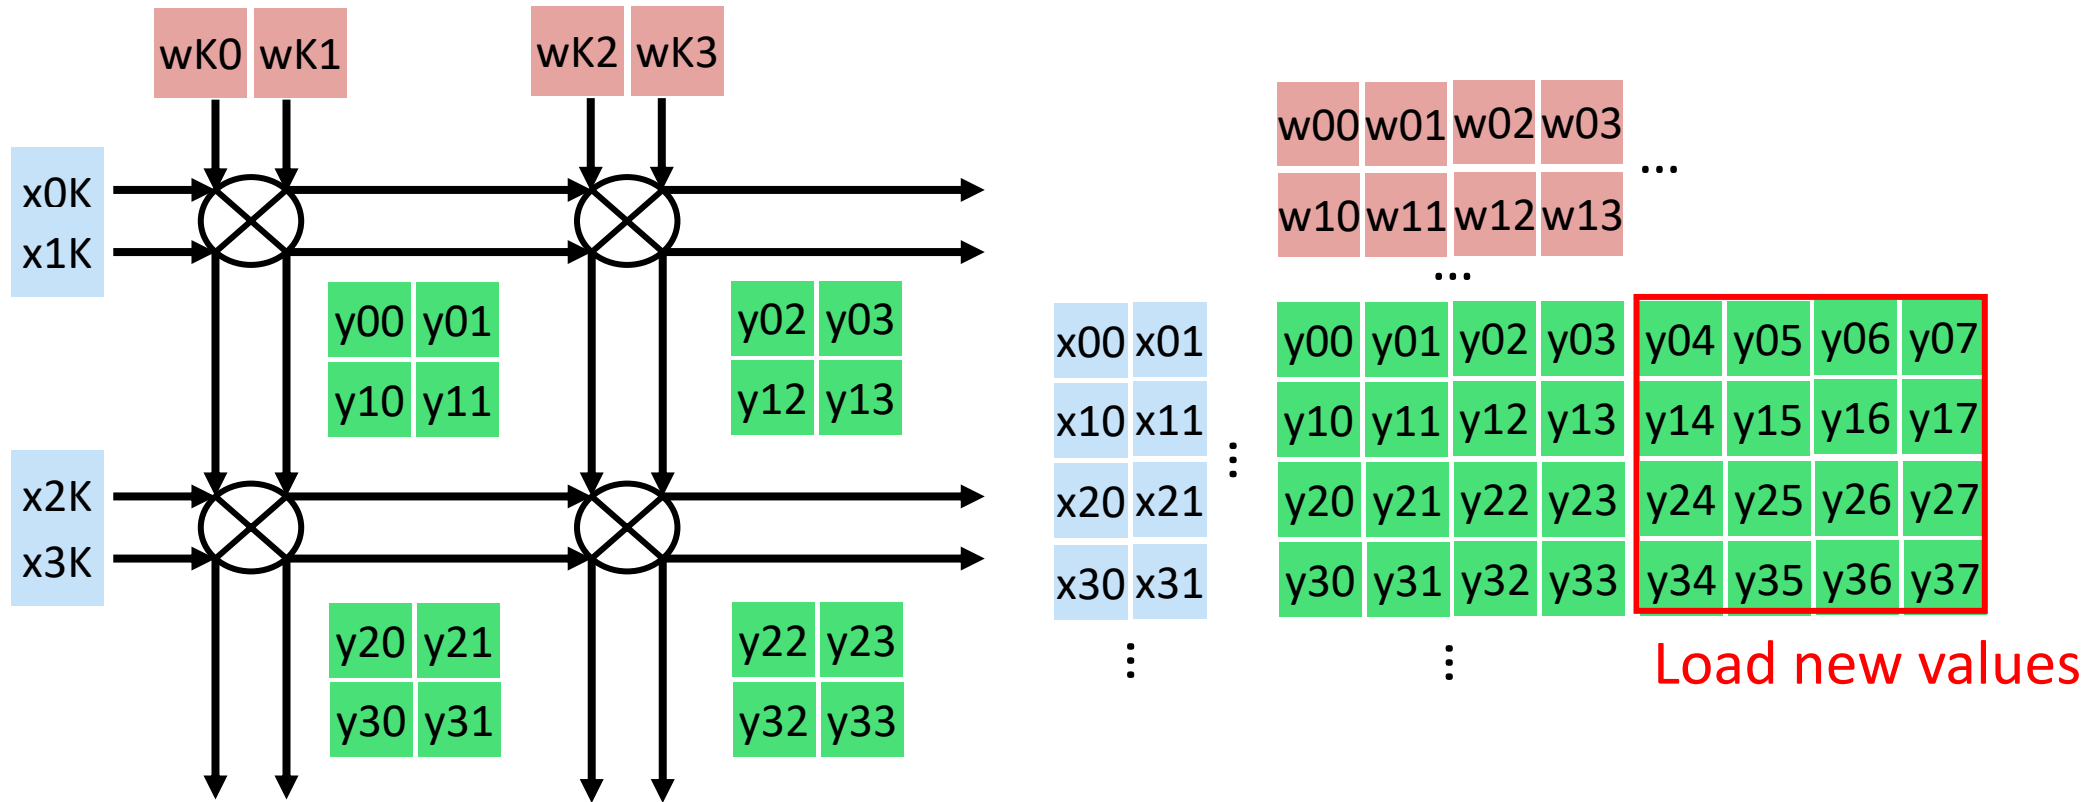
Output stationary engine, outer product



The FPU has 4 pipeline stages, to keep all of them busy we must load 2 elements from X and 2 elements from W per FPU every 4 cycles. Ys are stationary.



Output stationary engine, outer product



But once we arrive at the end of an X column / W row we need to replace them with a new output tile, we can exploit the 4 computation cycles, pre-loading 4ys/FPU every 4 cycles

Output stationary engine , outer product



Given L the size of the tensor unit (number of FPUs on one side)

- 1. Load $4L^2$ from Y to initiate the accumulators**
- 2. 4-cycles operation of the accelerator**
 1. Load $2L$ from X
 2. Load $2L$ from W
 3. Load $2L$ from Y (next iteration)
 4. Load $2L$ from Y (next iteration)
- 3. In $4L$ cycles the computation on the initial $4L^2$ Y elements is completed and the accelerator can continue on the $4L^2$ Y -elements loaded during point 2**

Output stationary engine



What is the required BW?

- Square tensor engine
- P = number of pipeline stages of the FPU
- L = number of FPUs per tensor engine side

RedMule:

$$BW = \frac{[16bit \times L \times (P+1)]}{cycle} \quad \frac{BW}{\pi} = \frac{16 \times (P+1)}{L} \frac{bit}{OP}$$


Output stationary:

$$BW = \frac{[16bit \times 2 \times L]}{cycle} \quad \frac{BW}{\pi} = \frac{16 \times 2}{L} \frac{bit}{OP}$$



Latency tolerance



Increase latency tolerance with more accumulators



2 elements /
4 cycles



4 elements /
16 cycles

Next Steps



- **TensorPool:**

- Optimize RedMule to L1 connection → modify TeraPool interconnect for TensorPool-5 config.
- TeraPool + RedMule (TensorPool) physical design in 7nm
- PPA on model microkernels and operators (combined RedMule&Cores)

- **System Performance:**

- Simulation speed is impaired by large design size → from RTL simulation to higher abstraction level, e.g. GVSoC
- TensorPool GVSoC model developement
- Data-Movement and end-end performance