

NextRAN-AI

ETH- Huawei Sweden

Marco Bertuletti

mbertuletti@iis.ee.ethz.ch

Yichao Zhang

yiczhang@iis.ee.ethz.ch

Mahdi Abdollahpour

mahdi.abdollahpout@unibo.it

Alessandro Vanelli-Coralli


avanelli@iis.ee.ethz.ch

Luca Benini

lbenini@iis.ee.ethz.ch

PULP Platform

Open Source Hardware, the way it should be!



@pulp_platform



pulp-platform.org



youtube.com/pulp_platform



We choosed to explore NeuralRX



Advantages of NeuralRX over other models

- **Flexible** = the same trained model supports different number of users, different number of subcarriers, different modulation schemes
- It generalizes well to many different channel models
- It is open-sourced and tested already on a real-time and standard compliant scenario (NeuralRX RT)

Next Steps:

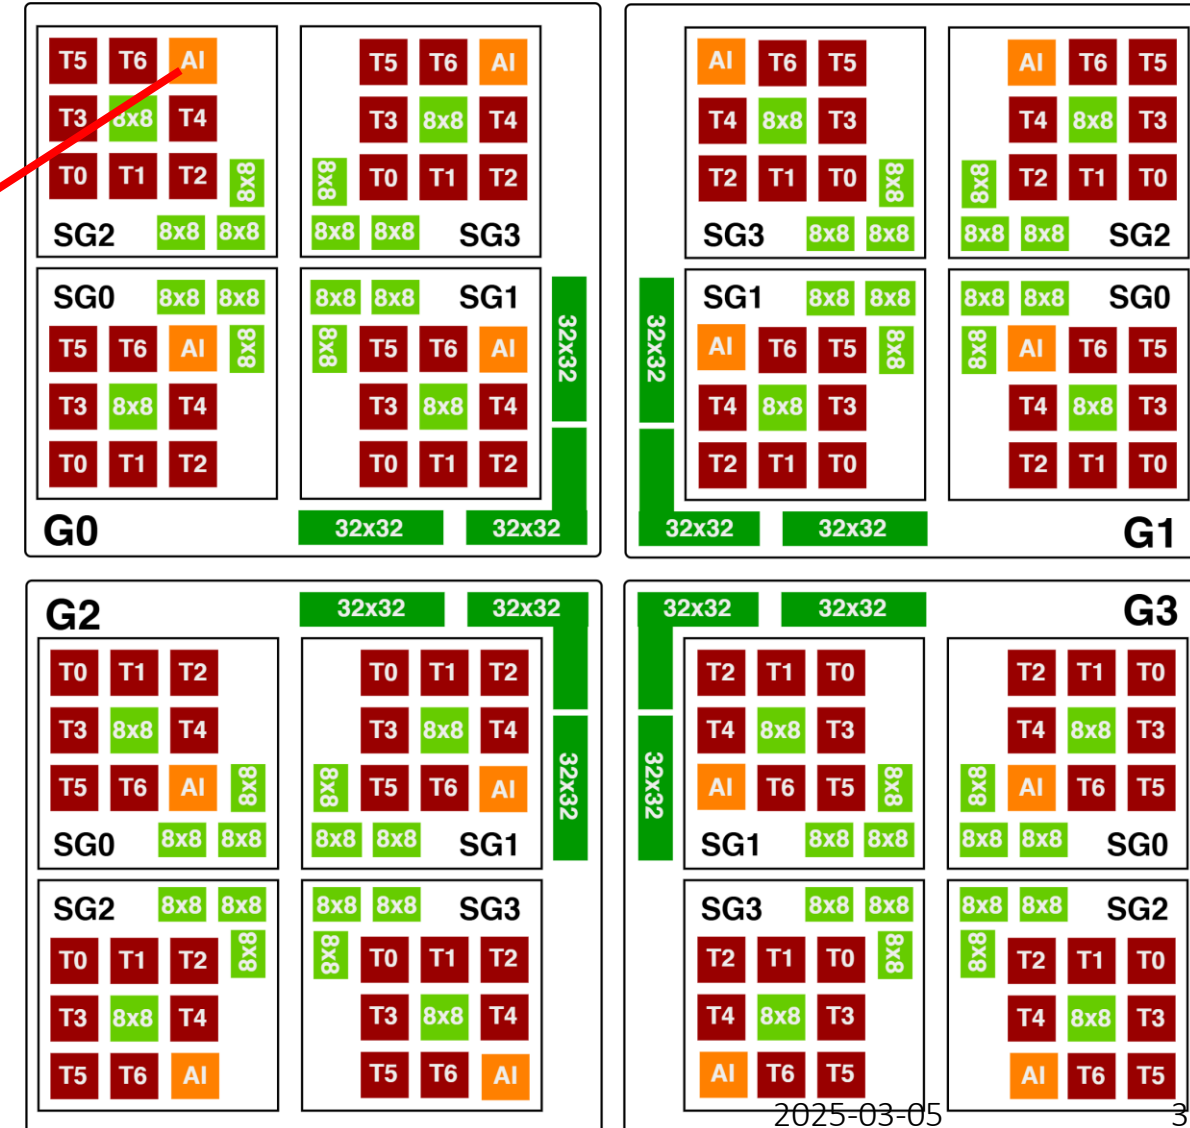
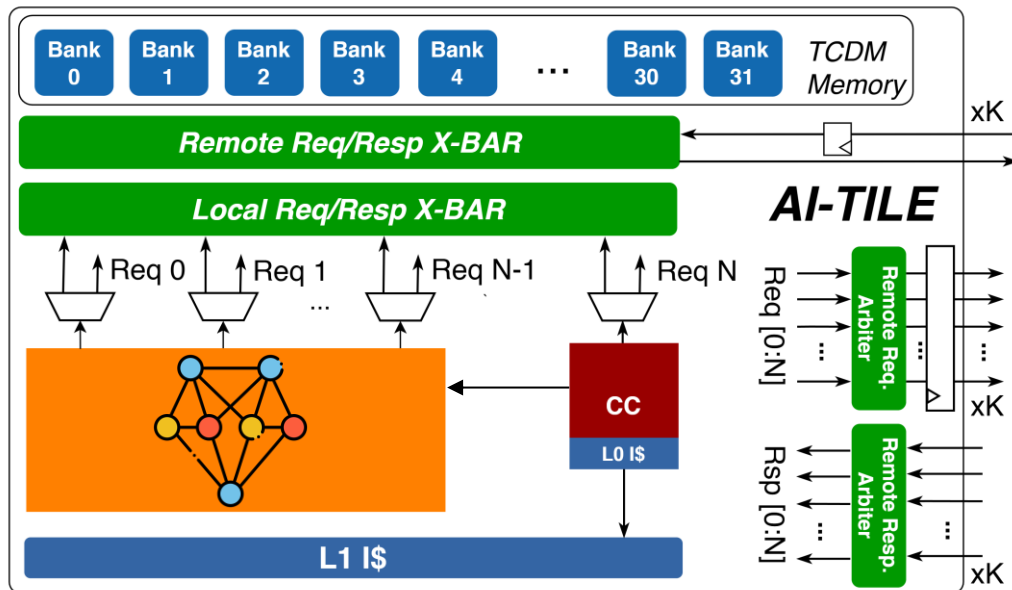
- Reduce model size and computational complexity for edge-deployment
- Possibly extend to more subcarriers, transceivers
- **Adequate TeraPool's computation per cycle**

We will present these results on next meeting

NextRAN-AI will increase TeraPool's compute capabilities



- Add AI-specialization (more compute), keeping configurability

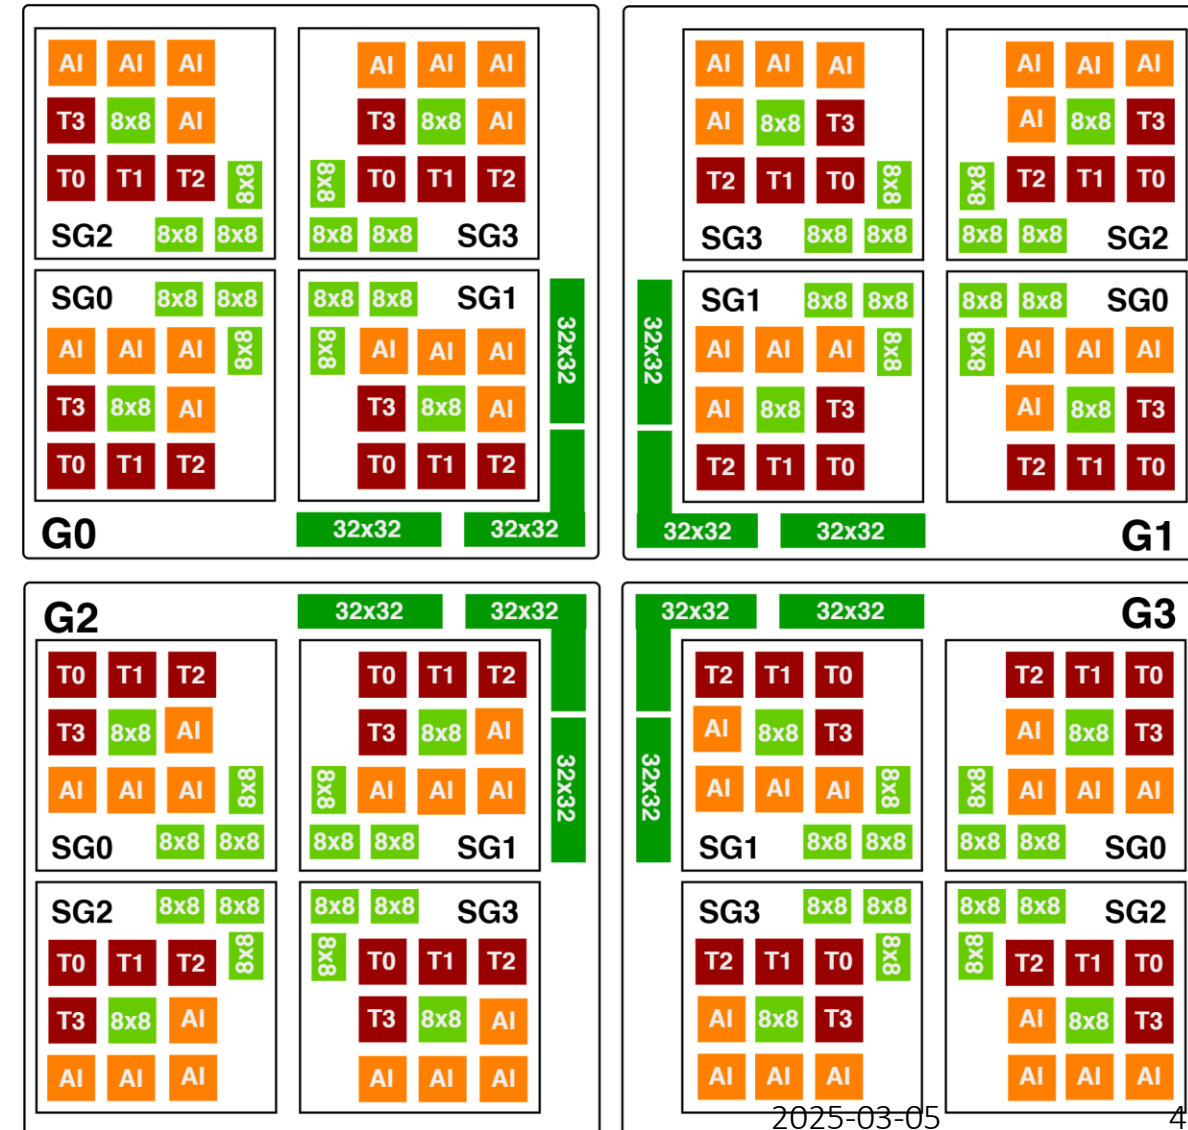


NextRAN-AI will increase TeraPool's compute capabilities



- Add AI-specialization (more compute), keeping configurability

How many accelerators?



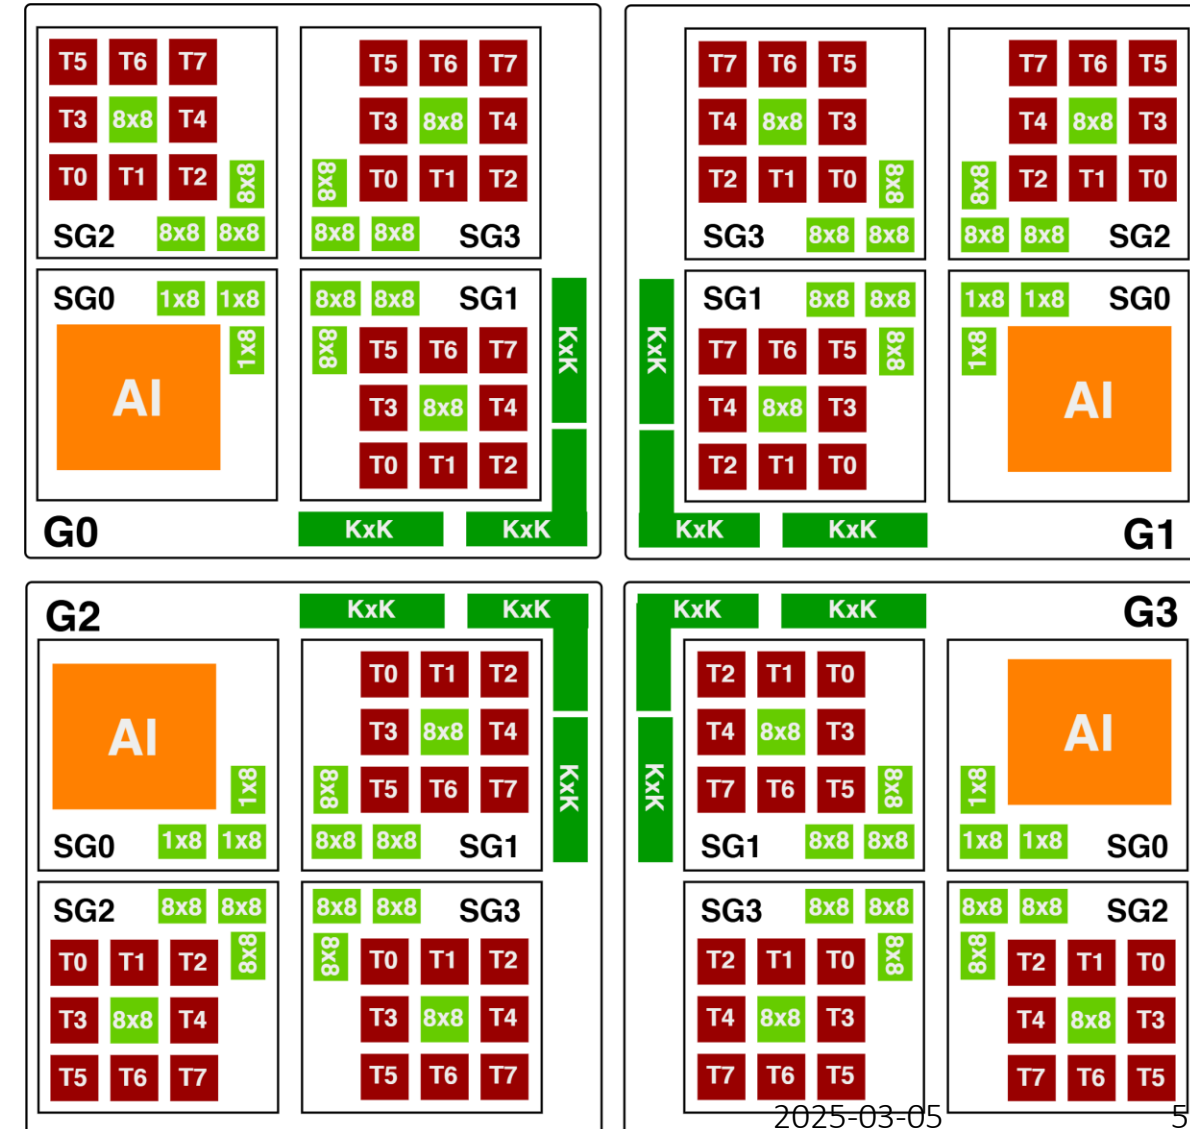
NextRAN-AI will increase TeraPool's compute capabilities



- Add AI-specialization (more compute), keeping configurability

How many accelerators?

How big?



NextRAN-AI will increase TeraPool's compute capabilities

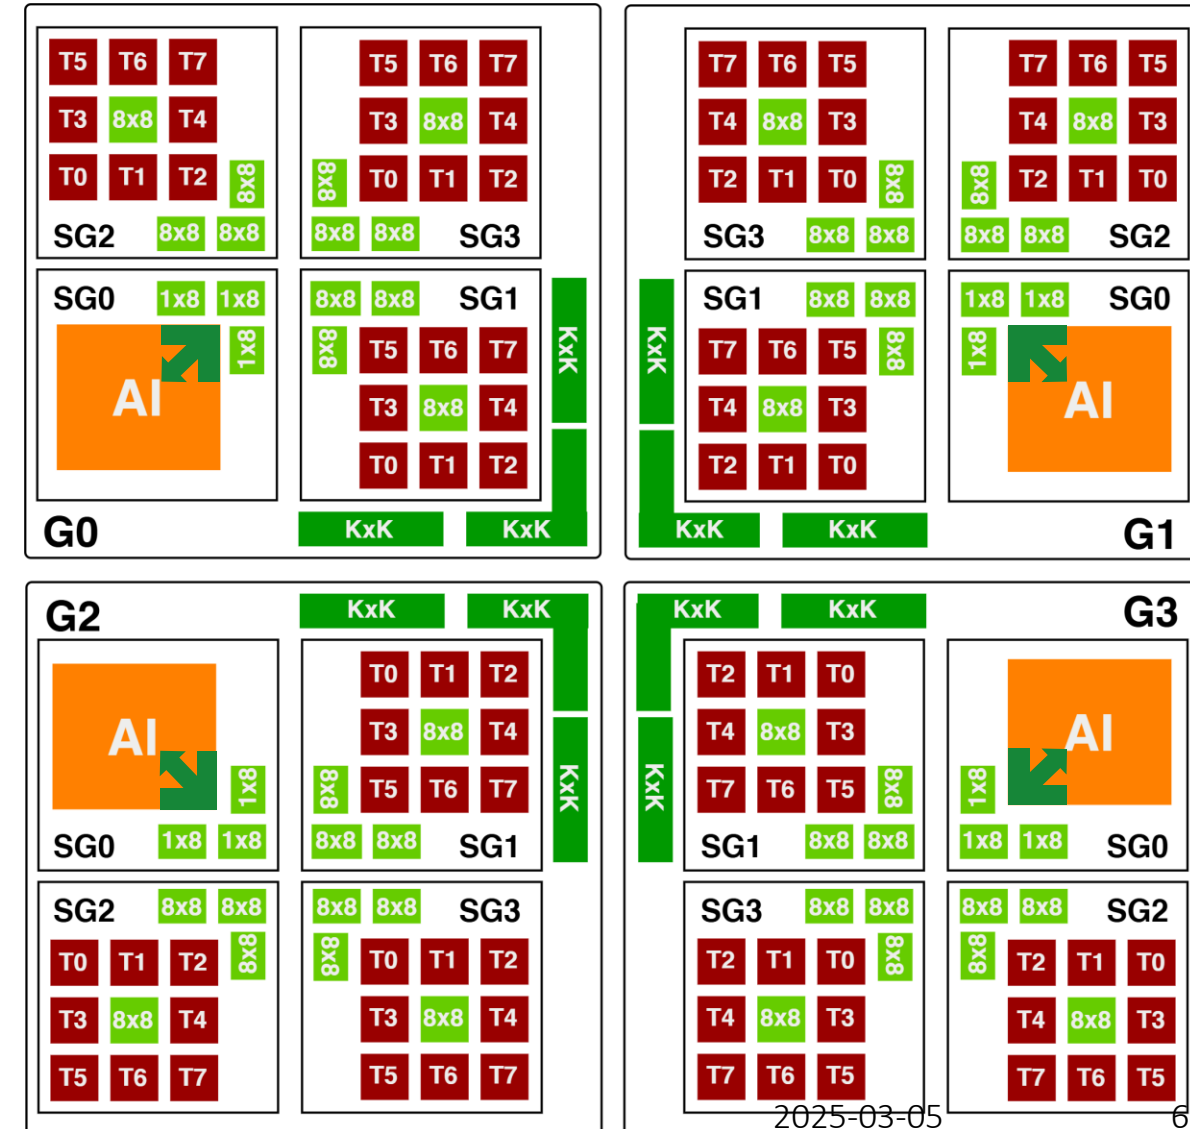


- Add AI-specialization (more compute), keeping configurability

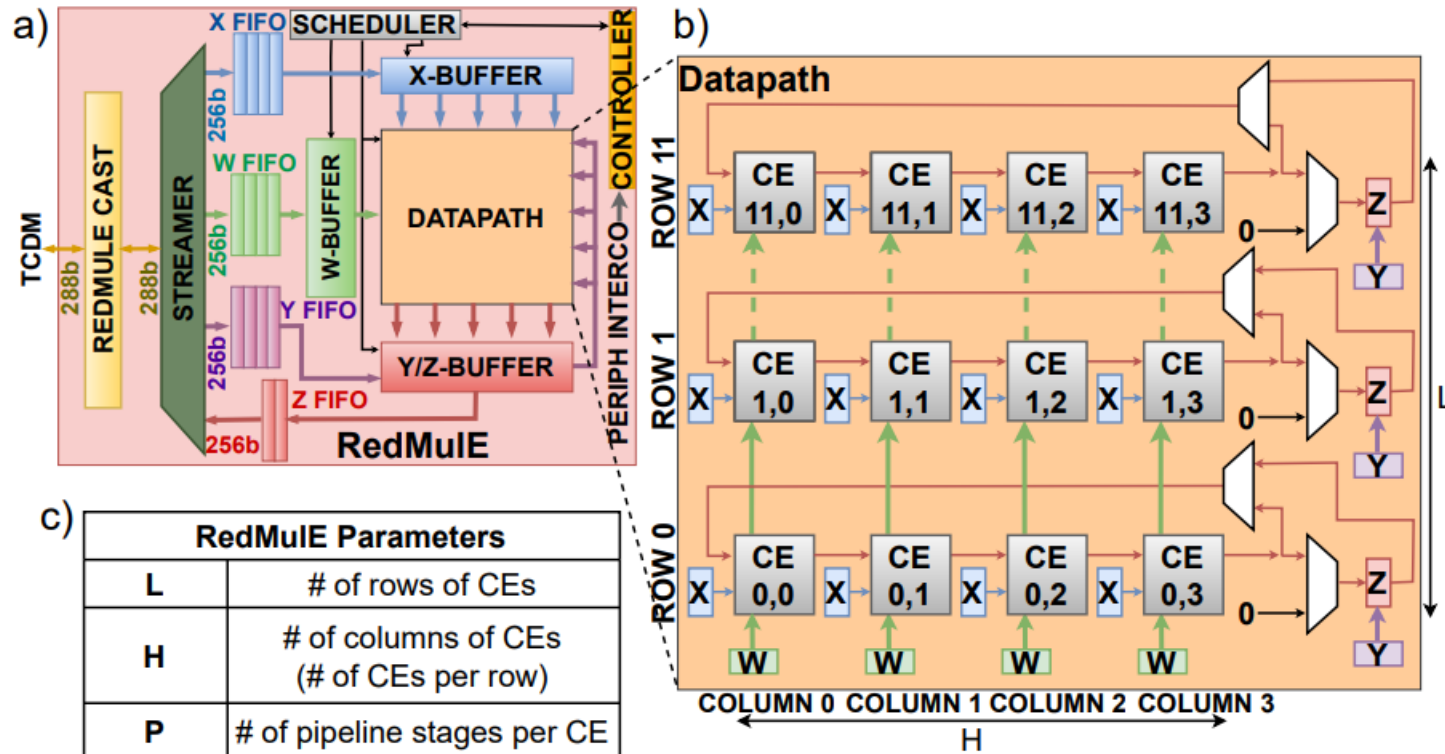
How many accelerators?

How big?

How connected?



RedMule - GEMM Accelerator



- **General Purpose** →
 - Attention/1D-Conv
 - Hermitian calculation
 - Beamforming
- **Open-Source**
- **Parametrizable**
- **TCDM-compatible**
- **Programmed via register interface**

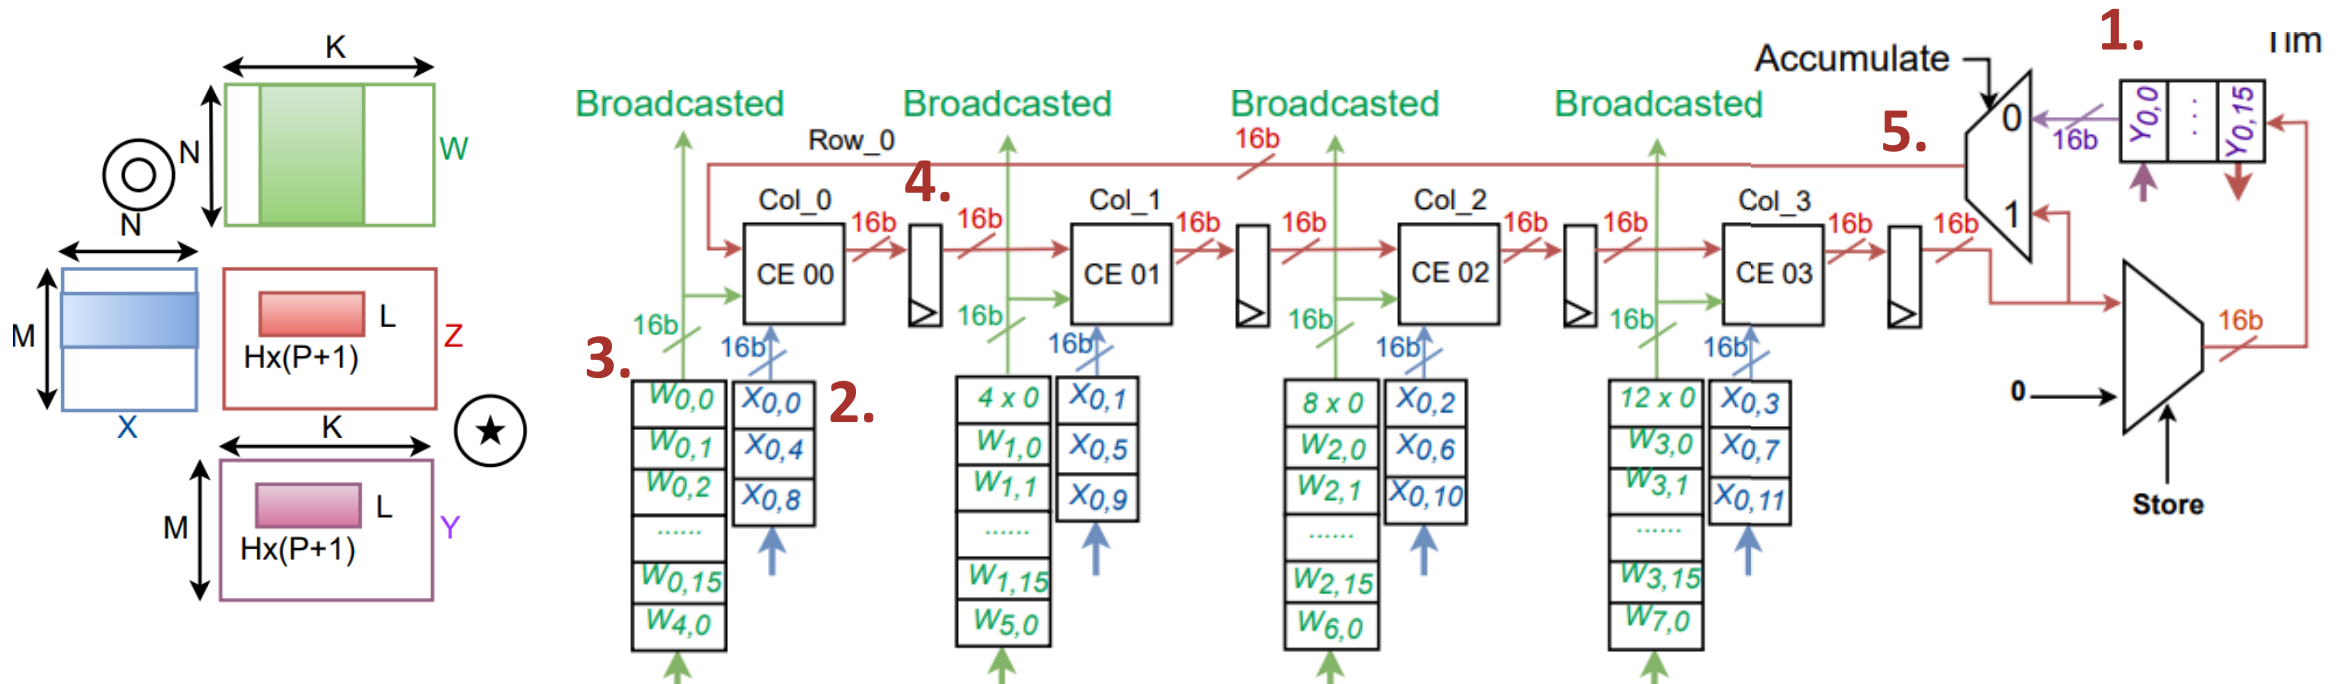
github.com/pulp-platform/snitch



RedMulE – Operation $Z = Y \star (X \odot W)$



1. Preload the Z buffer with L rows ($H \cdot (P+1) \cdot 16b$ elements each) from Y-matrix
2. Preload the X buffer with L rows ($H \cdot (P+1) \cdot 16b$ elements each) from X-matrix
3. Load $H \cdot (P+1)$ 16b weights from W-matrix, broadcast to all CE on a column
4. After $(P+1)$ cycles pass the result to next CE, load other $H \cdot (P+1)$ 16b weights
5. After $H \cdot (P+1)$ cycles **feedback**

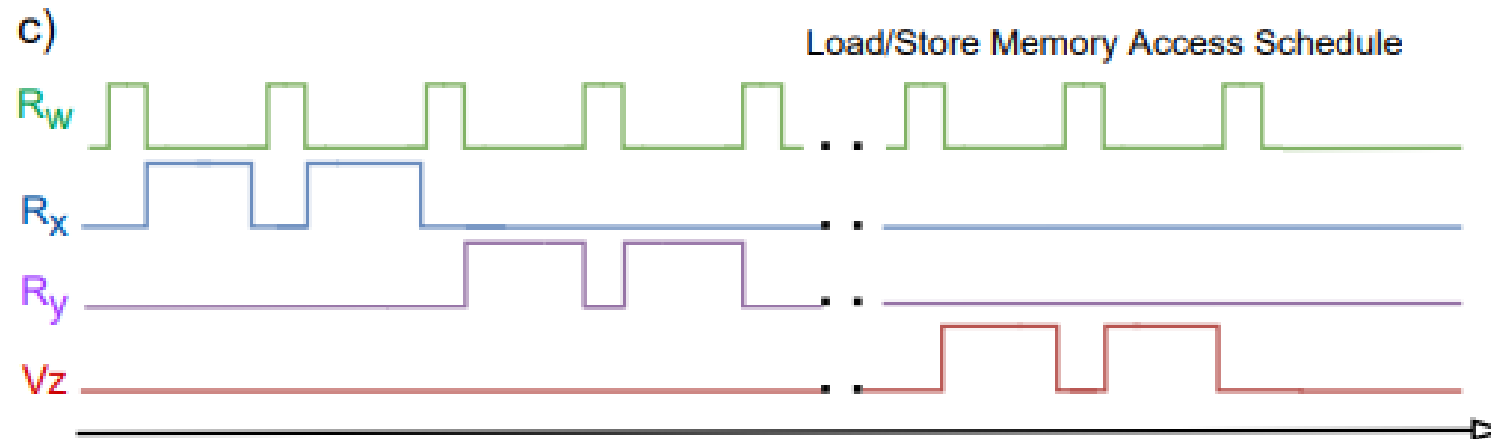


Designed for 1-cycle memory latency



- RedMuleE has a $16b * H * (P+1)$ memory port to TCDM
- The memory requests are split by the TCDM protocol DataWidth (32b)
- If the memory has one cycle latency \rightarrow Needs to be handled for TeraPool


Memory requests to different buffers are interleaved to access to $16b * H * (P+1)$ per cycle



Parametrization – How big accelerator?




1. **Assumption:** 1-cycle access to memory
2. **Design-Choice:** 1 RedMulE per Tile
3. **Assumption:** 2KiB-bank, 32-banks/Tile (memory occupation **50%** or 100%)

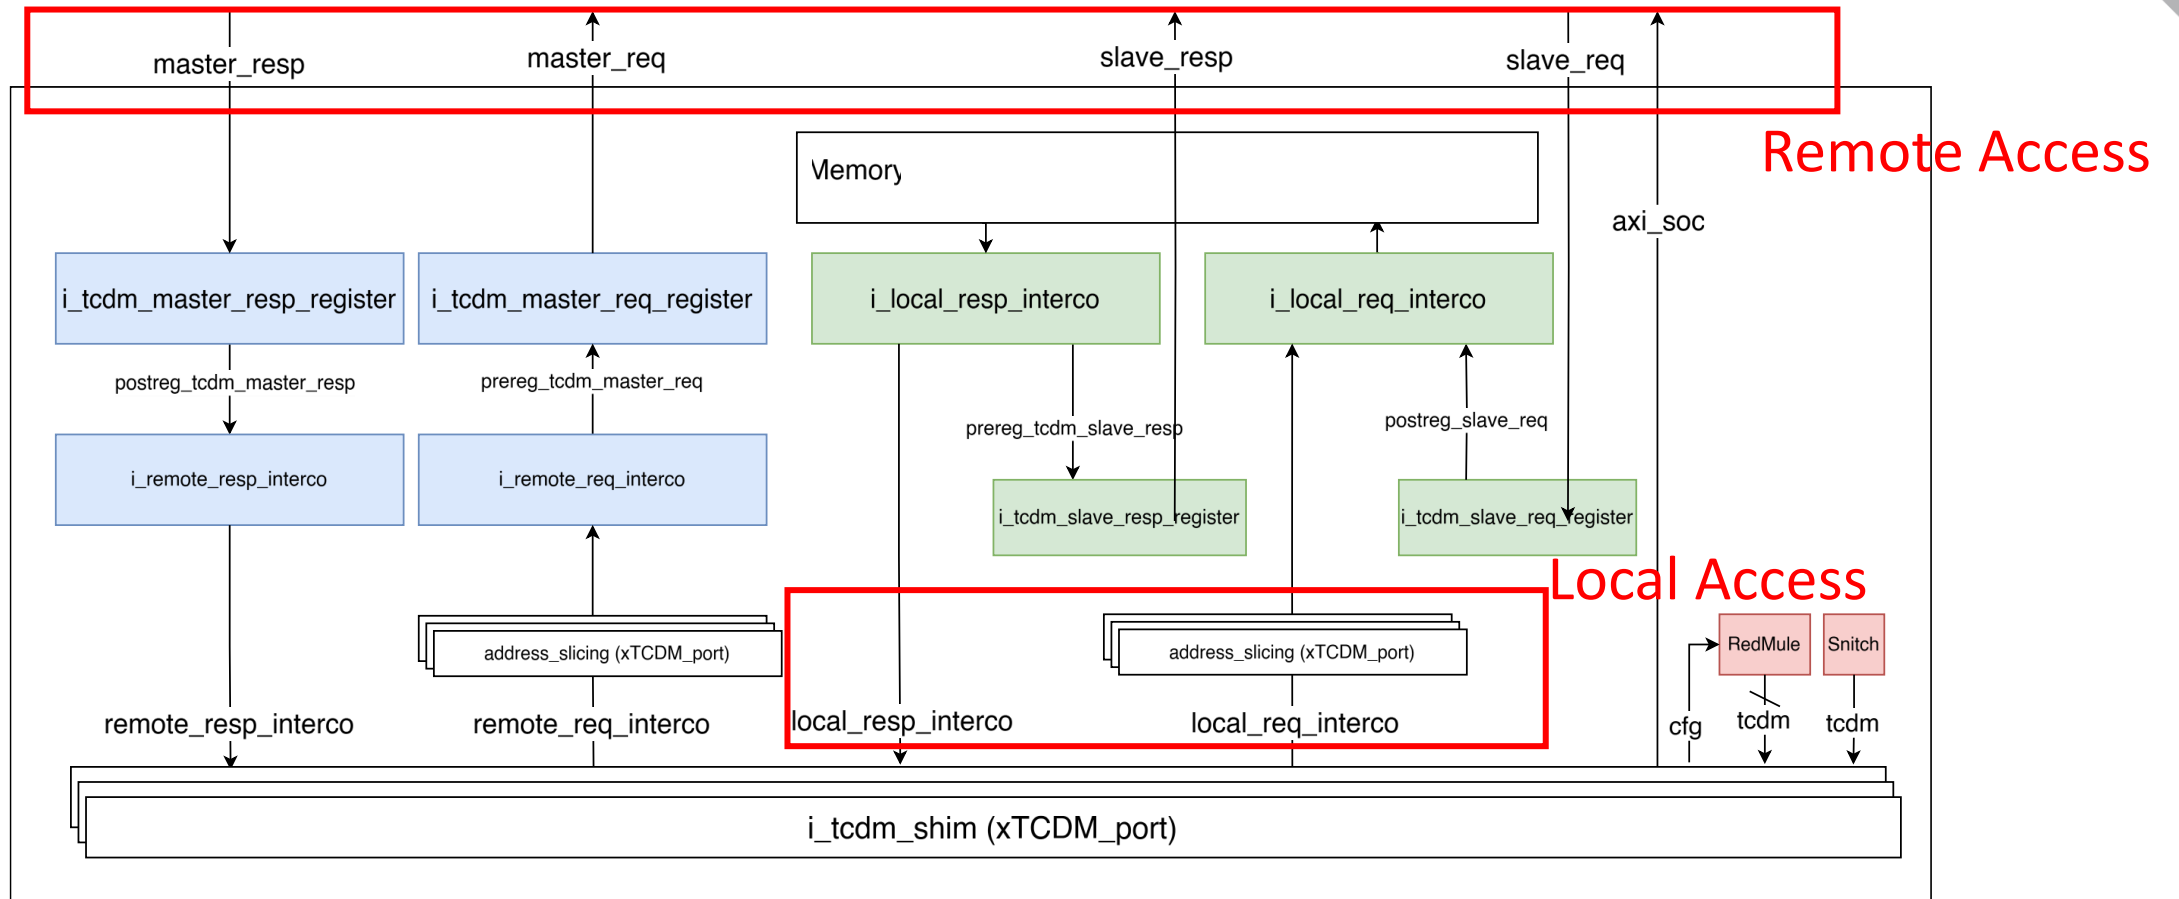


Double-buffering

Massive MIMO



Integration of RedMule in the Tile



- RedMule programmed by the Snitch core
- Parametrizable RedMule accesses TCDM (32b at a time through remote ports)

Next steps



Low utilization of the CE because of NUMA latency

- Add functionalities to restrict weights/inputs allocation to the **local memory** (inside a Tile we have 1-cycle access)
- For **remote access** improve latency on large requests (e.g. bursts)

